

```

# Required Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

# Step 1: Generate a Synthetic Dataset
np.random.seed(42)
data = {
    'battery_power': np.random.randint(500, 4000, 2000),
    'blue': np.random.randint(0, 2, 2000),
    'clock_speed': np.random.uniform(0.5, 3.0, 2000),
    'dual_sim': np.random.randint(0, 2, 2000),
    'fc': np.random.randint(0, 20, 2000),
    'four_g': np.random.randint(0, 2, 2000),
    'int_memory': np.random.randint(2, 128, 2000),
    'm_dep': np.random.uniform(0.1, 1.0, 2000),
    'mobile_wt': np.random.randint(80, 250, 2000),
    'n_cores': np.random.randint(1, 9, 2000),
    'pc': np.random.randint(0, 20, 2000),
    'px_height': np.random.randint(0, 1960, 2000),
    'px_width': np.random.randint(500, 2000, 2000),
    'ram': np.random.randint(256, 8192, 2000),
    'sc_h': np.random.randint(5, 19, 2000),
    'sc_w': np.random.randint(0, 18, 2000),
    'talk_time': np.random.randint(2, 20, 2000),
    'three_g': np.random.randint(0, 2, 2000),
    'touch_screen': np.random.randint(0, 2, 2000),
    'wifi': np.random.randint(0, 2, 2000),
    'price_range': np.random.randint(0, 4, 2000)
}

df = pd.DataFrame(data)

# Step 2: Read the Mobile Price Dataset
# Already generated above as 'df'

# Step 3: Print the First Five Rows
print("First five rows of the dataset:")
print(df.head())

# Step 4: Basic Statistical Computations
print("\nBasic statistical description of the dataset:")
print(df.describe())

# Step 5: Columns and Their Data Types
print("\nColumns and their data types:")
print(df.dtypes)

# Step 6: Detect and Handle Null Values
print("\nChecking for null values:")
print(df.isnull().sum())

# In case of null values, fill with mode (here it's simulated, so there are no nulls)
# for column in df.columns:
#     df[column].fillna(df[column].mode()[0], inplace=True)

# Step 7: Data Exploration Using Heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()

# Step 8: Split the Data into Train and Test Sets
X = df.drop('price_range', axis=1)
y = df['price_range']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 9: Fit the Naive Bayes Classifier
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)

# Step 10: Predict and Evaluate the Model

```

```
y_pred = nb_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print(f"\nAccuracy of the Naive Bayes model: {accuracy:.2f}")
```



First five rows of the dataset:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	\
0	3674	0	2.372066	0	10	0	57	
1	1360	0	2.127869	0	11	1	91	
2	1794	0	2.052321	0	4	0	96	
3	1630	1	1.380937	0	15	0	121	
4	1595	1	2.603619	1	12	0	52	

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	\
0	0.542688	96	7	...	727	1909	6217	18	11	
1	0.117649	183	5	...	627	693	8016	6	17	
2	0.704231	184	5	...	117	1626	2653	11	2	
3	0.570842	183	6	...	269	1899	3980	14	4	
4	0.908775	226	8	...	1418	786	7453	14	13	

	talk_time	three_g	touch_screen	wifi	price_range
0	4	0		1	1
1	5	0		1	0
2	19	1		1	1
3	15	0		0	3
4	16	0		1	2

[5 rows x 21 columns]

Basic statistical description of the dataset:

	battery_power	blue	clock_speed	dual_sim	fc	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	2282.193000	0.524000	1.738788	0.503500	9.435000	
std	1026.460304	0.499549	0.721614	0.500113	5.753288	
min	501.000000	0.000000	0.500029	0.000000	0.000000	
25%	1398.250000	0.000000	1.127818	0.000000	4.000000	
50%	2278.500000	1.000000	1.732105	1.000000	9.000000	
75%	3231.500000	1.000000	2.357399	1.000000	15.000000	
max	3999.000000	1.000000	2.998894	1.000000	19.000000	

	four_g	int_memory	m_dep	mobile_wt	n_cores	...	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	...	
mean	0.479500	64.818000	0.550504	163.210500	4.539500	...	
std	0.499705	36.576424	0.255575	49.845627	2.288899	...	
min	0.000000	2.000000	0.100048	80.000000	1.000000	...	
25%	0.000000	33.000000	0.332414	119.000000	3.000000	...	
50%	0.000000	65.000000	0.551112	162.000000	5.000000	...	
75%	1.000000	95.000000	0.767772	208.000000	7.000000	...	
max	1.000000	127.000000	0.999555	249.000000	8.000000	...	

	px_height	px_width	ram	sc_h	sc_w	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	979.965000	1263.289000	4280.538000	11.564000	8.463500	
std	567.443341	436.008547	2294.515552	4.125338	5.146932	
min	0.000000	502.000000	257.000000	5.000000	0.000000	
25%	482.000000	888.750000	2319.000000	8.000000	4.000000	
50%	977.000000	1266.000000	4253.500000	11.000000	8.000000	
75%	1472.000000	1635.000000	6319.750000	15.000000	13.000000	
max	1958.000000	1999.000000	8190.000000	18.000000	17.000000	

	talk_time	three_g	touch_screen	wifi	price_range
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	10.571500	0.475000	0.506000	0.493500	1.488500
std	5.171195	0.499499	0.500089	0.500083	1.111074
min	2.000000	0.000000	0.000000	0.000000	0.000000
25%	6.000000	0.000000	0.000000	0.000000	0.000000
50%	11.000000	0.000000	1.000000	0.000000	2.000000
75%	15.000000	1.000000	1.000000	1.000000	2.000000
max	19.000000	1.000000	1.000000	1.000000	3.000000

[8 rows x 21 columns]

Columns and their data types:

battery_power	int64
blue	int64
clock_speed	float64
dual_sim	int64
fc	int64
four_g	int64
int_memory	int64
m_dep	float64
mobile_wt	int64
n_cores	int64
pc	int64
px_height	int64
px_width	int64
ram	int64
sc_h	int64
sc_w	int64
...	...

```

talk_time      int64
three_g        int64
touch_screen   int64
wifi           int64
price_range    int64
dtype: object

```


Checking for null values:

```

battery_power  0
blue           0
clock_speed    0
dual_sim       0
fc            0
four_g         0
int_memory     0
m_dep          0
mobile_wt      0
n_cores        0
pc            0
px_height      0
px_width       0
ram            0
sc_h           0
sc_w           0
talk_time      0
three_g        0
touch_screen   0
wifi           0
price_range    0
dtype: int64

```

Feature Correlation Heatmap




```
# Step 1: Define the dataset
data = [
    ['Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same', 'Yes'],
    ['Sunny', 'Warm', 'High', 'Strong', 'Warm', 'Same', 'Yes'],
    ['Rainy', 'Cold', 'High', 'Strong', 'Warm', 'Change', 'No'],
    ['Sunny', 'Warm', 'High', 'Strong', 'Cool', 'Change', 'Yes']
]

# Step 2: Initialize the most specific hypothesis
hypothesis = ['0'] * (len(data[0]) - 1) # Exclude the target value ('Enjoy Sport')

# Step 3: Apply Find-S algorithm
for example in data:
    if example[-1] == 'Yes': # Only consider positive examples
        for i in range(len(hypothesis)):
            if hypothesis[i] == '0': # If the hypothesis is '0', set it to the attribute value
                hypothesis[i] = example[i]
            elif hypothesis[i] != example[i]: # If there's a conflict, set it to '?'
                hypothesis[i] = '?'

# Step 4: Print the final hypothesis
print("The most specific hypothesis is:")
print(hypothesis)
```

 The most specific hypothesis is:
['Sunny', 'Warm', '?', 'Strong', '?', '?']

```
# Required Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Step 1: Generate a Synthetic Dataset
np.random.seed(42)
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

# Convert to pandas DataFrame for easier manipulation
data = pd.DataFrame(data=np.hstack((X, y)), columns=["X", "y"])

# Step 2: Split the Dataset into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Fit a Linear Regression Model
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)


# Step 4: Predict the Target Values for the Test Set
y_pred = lin_reg.predict(X_test)

# Step 5: Evaluate the Model's Performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R^2 Score: {r2:.2f}")

# Plotting the results
plt.figure(figsize=(10, 6))
plt.scatter(X, y, color="blue", label="Data points")
plt.plot(X_test, y_pred, color="red", linewidth=2, label="Regression line")
plt.xlabel("X")
plt.ylabel("y")
plt.title("Linear Regression")
plt.legend()
plt.show()
```

Mean Squared Error: 0.65
R² Score: 0.81



```

# Required Libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

# Step 1: Generate a Synthetic Dataset
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=500, n_features=2, n_informative=2, n_redundant=0, n_clusters_per_class=1, random_state=42)

# Convert to pandas DataFrame for easier manipulation
data = pd.DataFrame(data=np.hstack((X, y.reshape(-1, 1))), columns=["Feature 1", "Feature 2", "Target"])

# Step 2: Split the Dataset into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Fit a KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Step 4: Predict the Target Values for the Test Set
y_pred = knn.predict(X_test)

# Step 5: Evaluate the Model's Performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Plotting the decision boundary
h = .02 # step size in the mesh
cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
cmap_bold = ListedColormap(['#FF0000', '#0000FF'])

# Create a mesh to plot the decision boundary
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# Predict class for each point in mesh
Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

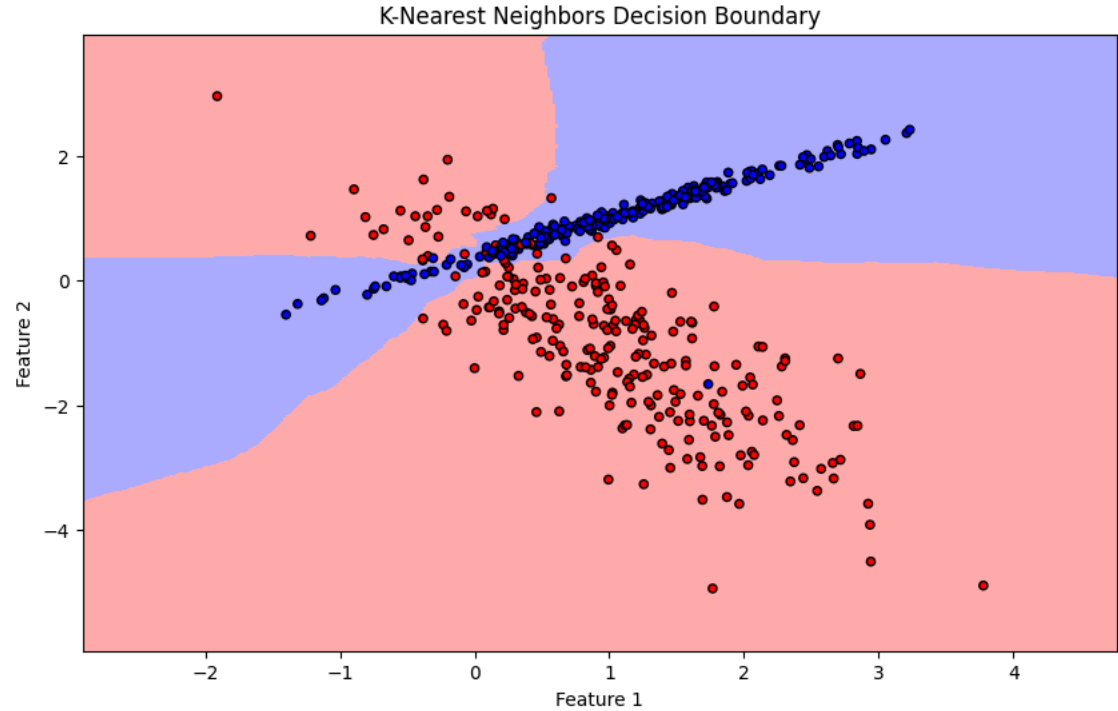
# Plot the decision boundary and training points
plt.figure(figsize=(10, 6))
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolor='k', s=20)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.title("K-Nearest Neighbors Decision Boundary")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()

```

↗

Accuracy: 0.97				
Classification Report:				
	precision	recall	f1-score	support
0	1.00	0.94	0.97	51
1	0.94	1.00	0.97	49
accuracy			0.97	100
macro avg	0.97	0.97	0.97	100
weighted avg	0.97	0.97	0.97	100

Confusion Matrix:
[[48 3]
[0 49]]




```

# Required Libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

# Step 1: Generate a Synthetic Dataset
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=500, n_features=2, n_informative=2, n_redundant=0, n_clusters_per_class=1, random_state=42)

# Convert to pandas DataFrame for easier manipulation
data = pd.DataFrame(data=np.hstack((X, y.reshape(-1, 1))), columns=["Feature 1", "Feature 2", "Target"])

# Step 2: Split the Dataset into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Fit a KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Step 4: Predict the Target Values for the Test Set
y_pred = knn.predict(X_test)

# Step 5: Evaluate the Model's Performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Plotting the decision boundary
h = .02 # step size in the mesh
cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
cmap_bold = ListedColormap(['#FF0000', '#0000FF'])

# Create a mesh to plot the decision boundary
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# Predict class for each point in mesh
Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot the decision boundary and training points
plt.figure(figsize=(10, 6))
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolor='k', s=20)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.title("K-Nearest Neighbors Decision Boundary")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()

```




Accuracy: 0.97

Classification Report:


	precision	recall	f1-score	support
0	1.00	0.94	0.97	51
1	0.94	1.00	0.97	49
accuracy			0.97	100
macro avg	0.97	0.97	0.97	100
weighted avg	0.97	0.97	0.97	100

Confusion Matrix:

```
[[48  3]
 [ 0 49]]
```



If you're already familiar with Colab, check out this video to learn about interactive tables, the executed code history view and the command palette.



```

import numpy as np
import pandas as pd

# Step 1: Define the dataset
data = [
    ['Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same', 'Yes'],
    ['Sunny', 'Warm', 'High', 'Strong', 'Warm', 'Same', 'Yes'],
    ['Rainy', 'Cold', 'High', 'Strong', 'Warm', 'Change', 'No'],
    ['Sunny', 'Warm', 'High', 'Strong', 'Cool', 'Change', 'Yes']
]

# Convert to pandas DataFrame
columns = ["Sky", "Air Temp", "Humidity", "Wind", "Water", "Forecast", "Enjoy Sport"]
df = pd.DataFrame(data, columns=columns)

# Step 2: Initialize the most specific hypothesis S0 and most general hypothesis G0
num_attributes = len(df.columns) - 1
S = ['0'] * num_attributes # Most specific hypothesis
G = [['?'] * num_attributes] # Most general hypothesis

# Step 3: Candidate Elimination Algorithm
def consistent(hypothesis, example):
    for i in range(len(hypothesis)):
        if hypothesis[i] != '?' and hypothesis[i] != example[i]:
            return False
    return True

for i, row in df.iterrows():
    example = row[:-1].tolist()
    target = row[-1]

    if target == 'Yes': # Positive example
        for j in range(len(S)):
            if S[j] == '0':
                S[j] = example[j]
            elif S[j] != example[j]:
                S[j] = '?'

        G = [g for g in G if consistent(g, example)]
    else: # Negative example
        G_temp = []
        for g in G:
            for j in range(num_attributes):
                if g[j] == '?':
                    for value in df.iloc[:, j].unique():
                        if value != example[j]:
                            new_g = g[:]
                            new_g[j] = value
                            if consistent(S, new_g):
                                G_temp.append(new_g)
            elif g[j] != example[j]:
                new_g = g[:]
                new_g[j] = '?'
                if consistent(S, new_g):
                    G_temp.append(new_g)
        G = G_temp

G = [g for g in G if any(consistent(g, row[:-1].tolist()) for _, row in df.iterrows())]

print("Final Specific Hypothesis S:")
print(S)
print("Final General Hypotheses G:")
for g in G:
    print(g)

🔗 Final Specific Hypothesis S:
['Sunny', 'Warm', '?', 'Strong', '?', '?']
Final General Hypotheses G:

```

```

# Required Libraries
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Step 1: Load the Dataset
iris = load_iris()
X = iris.data
y = iris.target

# Convert to binary classification: 1 if Iris-setosa, 0 otherwise
y = (y == 0).astype(int) # Setosa is class 0 in the original dataset

# Step 2: Preprocess the Data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)

# Step 3: Fit a Logistic Regression Model
log_reg = LogisticRegression(random_state=42)
log_reg.fit(X_train, y_train)

# Step 4: Predict the Target Values for the Test Set
y_pred = log_reg.predict(X_test)

# Step 5: Evaluate the Model's Performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Generate a classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Generate a confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```



Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	26
1	1.00	1.00	1.00	19
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Confusion Matrix:

```
[[26  0]
 [ 0 19]]
```

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture

# Generate synthetic data
np.random.seed(42)
n_samples = 500

# Generate random samples from two different normal distributions
shifted_gaussian = np.random.randn(n_samples, 2) + np.array([5, 5])
stretched_gaussian = np.dot(np.random.randn(n_samples, 2), np.array([[0.6, -0.6], [-0.4, 0.8]]))

X = np.vstack([shifted_gaussian, stretched_gaussian])

# Plot the synthetic data
plt.figure(figsize=(8, 6))
plt.scatter(X[:, 0], X[:, 1], s=5, alpha=0.5)
plt.title('Synthetic Data for GMM')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()

# Step 1: Initialize the Gaussian Mixture Model
gmm = GaussianMixture(n_components=2, max_iter=100, random_state=42)

# Step 2: Fit the model to the data
gmm.fit(X)

# Step 3: Predict the cluster for each data point
labels = gmm.predict(X)

# Step 4: Extract the parameters
weights = gmm.weights_
means = gmm.means_
covariances = gmm.covariances_

print("Weights:", weights)
print("Means:", means)
print("Covariances:", covariances)
```