

```

import time
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# Load the Iris dataset
iris = load_iris()
X, y = iris.data, iris.target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initialize models
models = {
    'Decision Tree': DecisionTreeClassifier(),
    'Logistic Regression': LogisticRegression(max_iter=200),
    'KNN': KNeighborsClassifier()
}

# Dictionary to store the results
results = {}

# Train and evaluate each model
for model_name, model in models.items():
    start_time = time.time()
    model.fit(X_train, y_train)
    training_time = time.time() - start_time

    start_time = time.time()
    y_pred = model.predict(X_test)
    prediction_time = time.time() - start_time

    accuracy = accuracy_score(y_test, y_pred)

    results[model_name] = {
        'accuracy': accuracy,
        'training_time': training_time,
        'prediction_time': prediction_time
    }

# Create a DataFrame to compare the results
results_df = pd.DataFrame(results).T
results_df = results_df[['accuracy', 'training_time', 'prediction_time']]

print(results_df)

```



	accuracy	training_time	prediction_time
Decision Tree	1.0	0.011393	0.000276
Logistic Regression	1.0	0.043086	0.000322
KNN	1.0	0.005454	0.036225

```

import pandas as pd

# Load the dataset
data = {
    'Example': [1, 2, 3, 4],
    'Citations': ['Some', 'Many', 'Many', 'Many'],
    'Size': ['Small', 'Big', 'Medium', 'Small'],
    'In Library': ['No', 'No', 'No', 'No'],
    'Price': ['Affordable', 'Expensive', 'Expensive', 'Affordable'],
    'Editions': ['Few', 'Many', 'Few', 'Many'],
    'Buy': ['No', 'Yes', 'Yes', 'Yes']
}
df = pd.DataFrame(data)
df.set_index('Example', inplace=True)

# Extract features and target
features = df.columns[:-1]
target = df.columns[-1]

# Initialize the most specific and most general hypotheses
def get_most_specific_hypothesis(num_features):
    return ['0'] * num_features

def get_most_general_hypothesis(num_features):
    return ['?'] * num_features

S = get_most_specific_hypothesis(len(features))
G = [get_most_general_hypothesis(len(features))]

# Function to check if one hypothesis is more general than another
def more_general(h1, h2):
    more_general_parts = []
    for x, y in zip(h1, h2):
        mg = x == '?' or (x != '0' and (x == y or y == '0'))
        more_general_parts.append(mg)
    return all(more_general_parts)

# Update S and G based on the training examples
for index, row in df.iterrows():
    if row[target] == 'Yes': # Positive example
        # Remove hypotheses from G that are inconsistent with the example
        G = [g for g in G if more_general(g, row[:-1].values)]
        # Update S to be consistent with the example
        for i, val in enumerate(row[:-1]):
            if S[i] == '0':
                S[i] = val
            elif S[i] != val:
                S[i] = '?'
        # Remove more general hypotheses from G
        G = [g for g in G if more_general(g, S)]
    else: # Negative example
        # Update G to be consistent with the example
        G_new = []
        for g in G:
            for i, val in enumerate(row[:-1]):
                if g[i] == '?':
                    for x in set(df[features[i]]):
                        if x != val:
                            new_hypothesis = g.copy()
                            new_hypothesis[i] = x
                            if more_general(new_hypothesis, S):
                                G_new.append(new_hypothesis)
            elif g[i] != val:
                G_new.append(g)
        G = G_new
        G = [g for g in G if more_general(g, S)]

print(f"Most Specific Hypothesis S: {S}")
print(f"Set of General Hypotheses G: {G}")

```

```

Most Specific Hypothesis S: ['Many', '?', 'No', '?', '?']
Set of General Hypotheses G: [['Many', '?', '?', '?', '?']]

```

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Generate a synthetic dataset
np.random.seed(0)
X = 2 - 3 * np.random.normal(0, 1, 100)
y = X - 2 * (X ** 2) + 1.5 * (X ** 3) + np.random.normal(-3, 3, 100)

X = X[:, np.newaxis]

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Transform features to polynomial features
poly = PolynomialFeatures(degree=3)
X_poly_train = poly.fit_transform(X_train)
X_poly_test = poly.transform(X_test)

# Train the polynomial regression model
model = LinearRegression()
model.fit(X_poly_train, y_train)

# Make predictions
y_train_pred = model.predict(X_poly_train)
y_test_pred = model.predict(X_poly_test)


# Evaluate the model
mse_train = mean_squared_error(y_train, y_train_pred)
r2_train = r2_score(y_train, y_train_pred)
mse_test = mean_squared_error(y_test, y_test_pred)
r2_test = r2_score(y_test, y_test_pred)

print(f"Training MSE: {mse_train}")
print(f"Training R^2: {r2_train}")
print(f"Test MSE: {mse_test}")
print(f"Test R^2: {r2_test}")

# Visualize the results
plt.scatter(X, y, color='blue', label='Data Points')
X_range = np.linspace(X.min(), X.max(), 100).reshape(-1, 1)
X_poly_range = poly.transform(X_range)
y_range_pred = model.predict(X_poly_range)
plt.plot(X_range, y_range_pred, color='red', label='Polynomial Regression Curve')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.title('Polynomial Regression')
plt.show()

```

Training MSE: 9.063856548037108
 Training R²: 0.9997112440964597
 Test MSE: 11.180742120161817
 Test R²: 0.9995350181165003



```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

# Load the Iris dataset
iris = load_iris()
X, y = iris.data, iris.target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Standardize the features (important for KNN)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train the KNN model
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)

# Make predictions
y_pred = knn.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Visualize the results (for the first two features)
def plot_decision_boundaries(X, y, model, title="Decision Boundaries"):
    h = 0.02 # step size in the mesh
    cmap_light = ListedColormap(['#FFAAAA', '#AAFFAA', '#AAAAFF'])
    cmap_bold = ListedColormap(['#FF0000', '#00FF00', '#0000FF'])

    # Create mesh grid
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))
```

```

plot_decision_boundaries(X_train, y_train, knn_vis, title="KNN Decision Boundaries (k=3)")

# Predict for each point in the mesh grid
Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure()
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)

# Plot the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolor='k', s=20)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.title(title)
plt.show()

# Only visualize using the first two features
X_vis = X_train[:, :2]
y_vis = y_train
knn_vis = KNeighborsClassifier(n_neighbors=3)
knn_vis.fit(X_vis, y_vis)

plot_decision_boundaries(X_vis, y_vis, knn_vis, title="KNN Decision Boundaries (k=3)")

```



Accuracy: 1.0

Classification Report:


	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Confusion Matrix:

```

[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]

```



```

import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelBinarizer
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

# Load the Iris dataset
iris = load_iris()
X, y = iris.data, iris.target

```

```

# For simplicity, let's reduce the problem to binary classification (e.g., Setosa vs. Non-Setosa)
# This is necessary because the Perceptron is a binary classifier
y_binary = (y != 0).astype(int) # Setosa is 0, Versicolor and Virginica are 1

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y_binary, test_size=0.3, random_state=42)

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train the Perceptron model
perceptron = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
perceptron.fit(X_train, y_train)

# Make predictions
y_pred_train = perceptron.predict(X_train)
y_pred_test = perceptron.predict(X_test)

# Evaluate the model
accuracy_train = accuracy_score(y_train, y_pred_train)
accuracy_test = accuracy_score(y_test, y_pred_test)
print(f"Training Accuracy: {accuracy_train}")
print(f"Test Accuracy: {accuracy_test}")
print("Classification Report (Test Set):")
print(classification_report(y_test, y_pred_test))
print("Confusion Matrix (Test Set):")
print(confusion_matrix(y_test, y_pred_test))

# Visualize the decision boundaries for the first two features (optional)
def plot_decision_boundaries(X, y, model, title="Decision Boundaries"):
    h = 0.02 # step size in the mesh
    cmap_light = ListedColormap(['#FFAAAA', '#AAFFAA'])
    cmap_bold = ListedColormap(['#FF0000', '#00FF00'])

    # Create mesh grid
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                        np.arange(y_min, y_max, h))

    # Predict for each point in the mesh grid
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.figure()
    plt.pcolormesh(xx, yy, Z, cmap=cmap_light)

    # Plot the training points
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolor='k', s=20)
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.title(title)
    plt.show()

# Only visualize using the first two features
X_vis = X_train[:, :2]
y_vis = y_train
perceptron_vis = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
perceptron_vis.fit(X_vis, y_vis)

plot_decision_boundaries(X_vis, y_vis, perceptron_vis, title="Perceptron Decision Boundaries (Binary Classification)")

# Note: For multi-class classification, a one-vs-rest strategy can be implemented.

```

```

Training Accuracy: 1.0
Test Accuracy: 1.0
Classification Report (Test Set):

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	26
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45


Confusion Matrix (Test Set):

```

[[19  0]
 [ 0 26]]

```

Perceptron Decision Boundaries (Binary Classification)



```

import pandas as pd

# Create the dataset
data = {
    'Example': [1, 2, 3, 4],
    'Citations': ['Some', 'Many', 'Many', 'Many'],
    'Size': ['Small', 'Big', 'Medium', 'Small'],
    'In Library': ['No', 'No', 'No', 'No'],
    'Price': ['Affordable', 'Expensive', 'Expensive', 'Affordable'],
    'Editions': ['Few', 'Many', 'Few', 'Many'],
    'Buy': ['No', 'Yes', 'Yes', 'Yes']
}

df = pd.DataFrame(data)

# Extract the features and labels
features = df.columns[1:-1]
X = df[features]
y = df['Buy']

# Initialize the most specific hypothesis
hypothesis = ['0'] * len(features)

# Find-S algorithm
for i, row in X.iterrows():
    if y[i] == 'Yes': # Only consider positive examples
        for j in range(len(hypothesis)):
            if hypothesis[j] == '0':
                hypothesis[j] = row[j]
            elif hypothesis[j] != row[j]:
                hypothesis[j] = '?'

print(f"Most specific hypothesis: {hypothesis}")

Most specific hypothesis: ['Many', '?', 'No', '?', '?']

```

```

import numpy as np
import pandas as pd
from sklearn.datasets import make_regression, make_classification
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, accuracy_score, classification_report
import matplotlib.pyplot as plt

# Generate a synthetic dataset for regression
X_reg, y_reg = make_regression(n_samples=100, n_features=1, noise=0.1, random_state=42)

# Generate a synthetic dataset for classification
X_clf, y_clf = make_classification(n_samples=100, n_features=2, n_informative=2, n_redundant=0, n_classes=2, random_state=42)

# Split the regression data into training and test sets
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X_reg, y_reg, test_size=0.3, random_state=42)

# Split the classification data into training and test sets
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(X_clf, y_clf, test_size=0.3, random_state=42)

# Train Linear Regression model
linear_reg = LinearRegression()
linear_reg.fit(X_train_reg, y_train_reg)
y_pred_reg = linear_reg.predict(X_test_reg)

# Evaluate Linear Regression model
mse_reg = mean_squared_error(y_test_reg, y_pred_reg)
print(f"Linear Regression MSE: {mse_reg}")

# Train Logistic Regression model
logistic_reg = LogisticRegression()
logistic_reg.fit(X_train_clf, y_train_clf)
y_pred_clf = logistic_reg.predict(X_test_clf)

# Evaluate Logistic Regression model
accuracy_clf = accuracy_score(y_test_clf, y_pred_clf)
print(f"Logistic Regression Accuracy: {accuracy_clf}")
print("Classification Report:")
print(classification_report(y_test_clf, y_pred_clf))

# Visualize the results for Linear Regression
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.scatter(X_test_reg, y_test_reg, color='blue', label='Actual')
plt.plot(X_test_reg, y_pred_reg, color='red', label='Predicted')
plt.title('Linear Regression')
plt.xlabel('Feature')
plt.ylabel('Target')
plt.legend()

# Visualize the results for Logistic Regression
plt.subplot(1, 2, 2)
plt.scatter(X_test_clf[:, 0], X_test_clf[:, 1], c=y_test_clf, cmap='viridis', marker='o', edgecolor='k', s=100, label='Actual')
plt.scatter(X_test_clf[:, 0], X_test_clf[:, 1], c=y_pred_clf, cmap='coolwarm', marker='x', edgecolor='k', s=100, label='Predicted')
plt.title('Logistic Regression')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()

plt.show()

```





Linear Regression MSE: 0.010347302683438472

Logistic Regression Accuracy: 0.9666666666666667

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.94	0.97	16
1	0.93	1.00	0.97	14
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

<ipython-input-9-257a282f56c5>:55: UserWarning: You passed a edgecolor/edgecolors ('k')
 plt.scatter(X_test_clf[:, 0], X_test_clf[:, 1], c=y_pred_clf, cmap='coolwarm', marker



```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm

# Generate synthetic data from a mixture of two Gaussian distributions
np.random.seed(0)

# Parameters of the two Gaussian distributions
mean_true = [-2, 3]
variance_true = [1, 2]
weights_true = [0.6, 0.4]

# Generate data points
def generate_data(n_samples):
    data = []
    for i in range(n_samples):
        z = np.random.choice([0, 1], p=weights_true)
        data.append(np.random.normal(loc=mean_true[z], scale=np.sqrt(variance_true[z])))
    return np.array(data)

# Number of data points
n_samples = 1000
X = generate_data(n_samples)

# Plot the generated data
plt.figure(figsize=(8, 6))
plt.hist(X, bins=30, density=True, alpha=0.5, color='blue', label='Histogram of Data')

# Define the EM algorithm for Gaussian Mixture Models (GMMs)
def expectation_maximization(X, n_components, max_iter=100, tol=1e-4):
    n_samples = len(X)
    # Initialize parameters randomly
    np.random.seed(0)
    means = np.random.randn(n_components)
    variances = np.random.rand(n_components)
    weights = np.random.rand(n_components)
```

```
weights /= np.sum(weights)

# Initialize old parameters for convergence check
means_old = np.zeros_like(means)
variances_old = np.zeros_like(variances)

# EM algorithm
for iter in range(max_iter):
    # E step: compute responsibilities
    responsibilities = np.zeros((n_samples, n_components))
    for k in range(n_components):
        responsibilities[:, k] = weights[k] * norm.pdf(X, loc=means[k], scale=np.sqrt(variances[k]))
    responsibilities /= responsibilities.sum(axis=1, keepdims=True)

    # M step: update parameters
    Nk = responsibilities.sum(axis=0)
    means = (responsibilities.T @ X) / Nk
    variances = (responsibilities.T @ (X**2) / Nk) - means**2
    weights = Nk / n_samples

    # Check for convergence
    if np.max(np.abs(means - means_old)) < tol and np.max(np.abs(variances - variances_old)) < tol:
        break

    means_old = means.copy()
    variances_old = variances.copy()
```